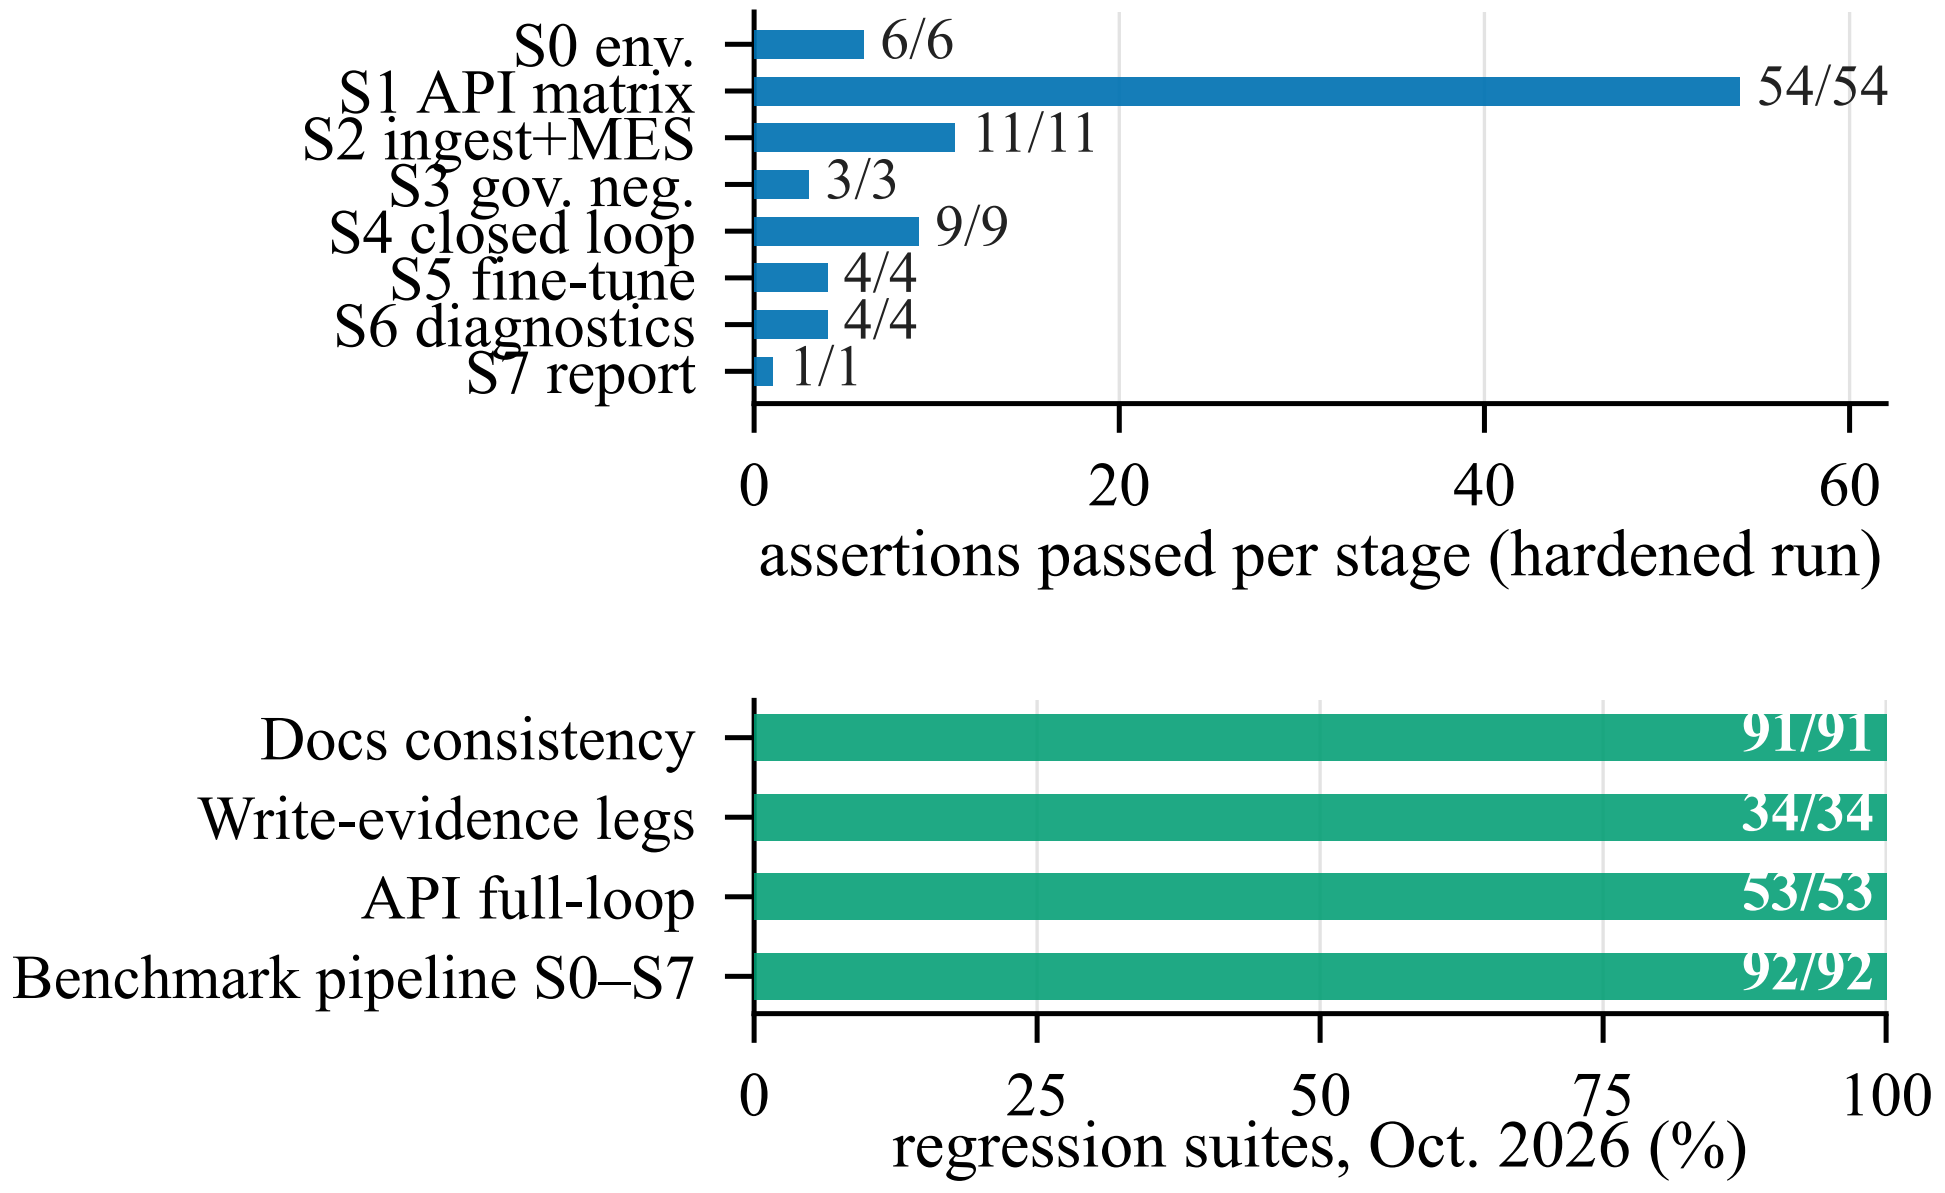
S0 env.
S1 API matrix
S2 ingest+MES
S3 gov. neg.
S4 closed loop
S5 fine-tune
S6 diagnostics
S7 report
6/6
54/54
11/11
3/3
9/9
4/4
4/4
1/1
0
20
40
60
assertions passed per stage (hardened run)
Docs consistency
Write-evidence legs
API full-loop
Benchmark pipeline S0–S7
91/91
34/34
53/53
92/92
0
25
50
75
100
regression suites, Oct. 2026 (%)